

Predicting Milwaukee Property Assessment Appeals

Quang Nguyen Adam Sperber Anthony Truckenbrod

Loyola University Chicago

Statistical Consulting - Fall 2021

Abstract

Rising housing markets and the COVID-19 pandemic led to record-high appeals of property assessments in 2020. In this study, we explored trends in which property owners appealed their assessments and methods for predicting which properties will appeal in the future. In particular, we investigated different classification methods for predicting property assessment appeal for residential buildings in Milwaukee. We focused on two supervised machine learning techniques: penalized logistic regression with a LASSO penalty and eXtreme Gradient Boosting (XGBoost). XGBoost was found to be the preferable model, as it performed the best in terms of AUC. Furthermore, analysis of feature importance indicated that location, year built, and finished area are the most important variables in predicting property assessment appeal for houses in the city of Milwaukee.

1 Introduction

The year 2020 saw a massive surge in housing prices nationwide, prompted by an already volatile market and exacerbated by a quarantined lifestyle which instilled many Americans' need for more physical space in their immediate surroundings. For many in Milwaukee who wished neither to sell nor buy, the increased prices served as little more than an unwanted increase in semi-annual property taxes. As a result, the mass influx of housing appeals alongside the housing market blurred the environment for, in the eyes of homeowners, obtaining an accurate property valuation.

The City of Milwaukee Assessor's Office (referred to herein as the Office) bears the responsibility of providing accurate valuations of properties throughout their respective cities. Over the past few years, this task of appraisal became much more difficult as values fluctuated quite dramatically over relatively short periods of time. Along with this responsibility, the Office manages and processes property assessment appeals, occurring when a property owner disagrees with their property's assessment for some specific reason; owners are then able to challenge the assessment by submitting a formal appeal. The Office then looks into the issue and determines whether or not to justify the appeal, acting accordingly to provide an accurate valuation of the property. Unsurprisingly, the rate of these property assessment appeals generally increases during times of market volatility; in Milwaukee, 2020 saw record rates of property assessment appeals, increasing fivefold over the previous year. This imbalance created an almost unmanageable additional workload for the Office during an already low-staff environment as a result of the pandemic. As a result, the concept of predicting which properties are likely to appeal their respective assessments became immensely valuable, as it would allow the Office to direct their resources to specific areas identified as more likely to appeal their property assessments.

In this report, we primarily investigated and identified trends of properties likely to appeal their assessments. Furthermore, we developed methods for predicting which properties are

likely to appeal said assessments in the future. These models will be used by the Office to improve initial valuations by identifying which properties/areas to review prior to giving assessments. The models will also be used to target public outreach to help educate taxpayers about the property assessment process given how many assessment appeals stem from a lack of understanding of property taxes.

The paper is outlined as follows. We first describe the data and perform exploratory analysis in Section 2. We then present our modeling strategy, methodology, and results in Section 3. Lastly, in Section 4, we give a summary of our main findings and discuss possible future work related to this project.

2 Data

The main dataset for our investigation consisted of 122,096 observations representing the properties throughout Milwaukee. Each property contained the following attributes: a unique property ID number, building type, the appraiser responsible for property assessment, and the quality (ranges from AA+ to E), and the condition of the building (ranges from excellent to unsound). In addition, the data included variables on the number of kitchens, full and half bathrooms alongside their respective rating (also from excellent to unsound), construction year, total finished area, land area, and the ultimate sale date and price. Finally, three columns of binary variables indicated whether the property owners appealed their appraisal during the tax years of 2019, 2020, and 2021.

Additionally, the Office provided us with independent property location data, each case containing an identification number, address, zip code, and census tract. Based on the given information, we performed geocoding and obtained the latitude and longitude for each property, using the `tidygeocoder` package (Cambon et al., 2021) in R (R Core Team, 2021). We merged the ensuing location table with the main dataset using the property ID as the

joining key.

After combining the datasets described above, we created visual and numerical summaries to explore our data. We first summarized the appeals looking for trends over time, as well as trends among the categorical and quantitative predictors. Table 1 shows the number of appeals and the appeal rates for 2019, 2020, and 2021. It is obvious that the proportion of property appeals in 2020 exceeded those from 2019 and 2021 by approximately five and nine times, respectively. Due to this imbalance, we and our client agreed to use the 2020 appeal rate as our final response variable as opposed to 2021, as it would allow for significantly higher model flexibility and performance.

Table 1: Number of appeals processed from 2019-2021 by the City of Milwaukee Assessor’s Office

Year	Number of appeals	% appealed
2019	886	0.726
2020	4392	3.597
2021	517	0.423

We then examined the proportion of missing data for our main housing dataset (Figure 1). We encountered that the property sale dates and prices were more than 80% missing. However, the data were not missing in a traditional sense; rather, sales were only recorded if they occurred in 2018 or later. The same could be said for the half-bathroom columns, count and rating, since many houses lacked them in the first place.



Figure 1: Proportion of missing data for each variable in the housing data. The variables are also ordered by the amount of missingness.

Next, we looked at the relationship between appeal status and building type. Initial comparisons between housing styles and appeal rate over the three years from 2019 to 2021 suggested a notable correlation between appeal rates and styles conventionally associated with more affluent lifestyles, most notably contemporary, mansion, and to a lesser extent Tudor-style houses. The appeal rates for these luxury residences normally top out around 4.5%, though 2020 exceeded 20% in contemporary-style buildings alone, the three aforementioned categories combining for well over 40% of all appeals.




Figure 2: Property appealed rate for all building types, broken down by year.

With respect to physical geography, Figure 3 displays the top zip codes in terms of appeal rate. These zip codes fall into two major location categories, namely, by the Lake Michigan Shoreline (including downtown), or situated away from the city near major commuter highways, most notably interstates 94 and 41. In addition, Figure 4 is a multi-panel plot of maps showing appeal rates across Milwaukee for the years between 2019 and 2021. This visualization aids the interpretability of our geographic data, though restricted to the city of Milwaukee. The highest-appealed zip codes, 53212, 53211, and 53207, all fall near downtown and the Lake Michigan Shore. All in all, property appeal is undoubtedly tied to geography, thus location variables are likely to be useful in predicting whether a homeowner appeals their property assessment.




Figure 3: Top 15 zip codes in Milwaukee with highest property appeal rate for each year between 2019 and 2021.

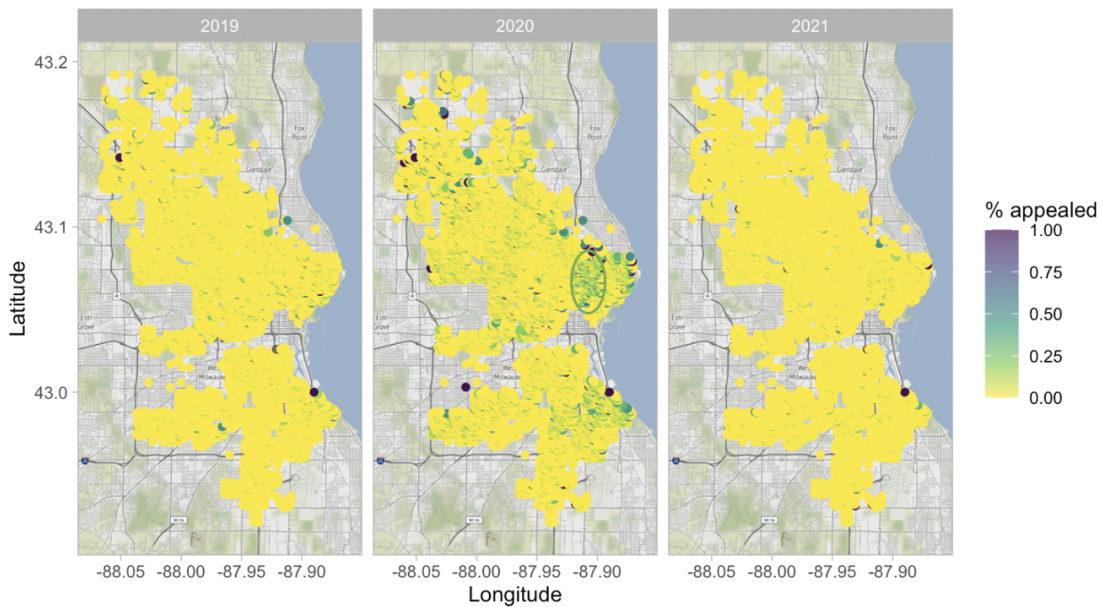


Figure 4: Map displaying appeal rates across Milwaukee for 2019, 2020, and 2021. The points representing the residential properties are color-coded by appeal rate, where darker colors indicate properties with higher appeal rates. Zip code 53212 is outlined in red in the central map.

3 Predictive Models

In this section, we present our modeling strategy for predicting property assessment appeals for residential properties in the city of Milwaukee. We explored the performance of different supervised machine learning approaches on the classification of our target variable, the appeal status for the tax year 2020.

3.1 Data Preparation

We started the modeling stage by doing feature engineering to transform the features in our raw data, making these features more useful in the machine learning process. For the ordinal categorical variables, we performed ordinal encoding to the columns representing rating with range poor to excellent. We converted the levels of these ordinal categorical predictors into integer values, with 1 representing the lowest level, poor. In addition, we converted the property quality grade column (containing values from AA+ to E-) to numeric scores, as suggested by our client. These are the weight values associated with each quality grade used by the assessor's office.

In addition, for each nominal categorical covariate, we implemented a one-hot encoding step. One-hot encoding transforms the predictors into columns of binary indicators (0 and 1), and for each variable, the number of new columns should equal the number of the levels belonging to the initial factor. Moreover, some of the variables in our dataset such as neighborhood, zip code, and building type contain rare categories (in one extreme case, a zip code partially outside the city of Milwaukee contained one appeal out of four total properties, overinflating its significance), so we decided to lump together the infrequently occurring levels in each of those factors into another level. Finally, we considered a natural log transformation for the continuous variables finished area and land in square feet.

3.2 Models

The first model we considered was a LASSO penalized logistic regression model (Tibshirani, 1996), which is short for Least Absolute Shrinkage and Selection Operator. LASSO involves penalizing the absolute magnitude of the regression coefficients, which is commonly known as the ℓ_1 penalty. For our data, LASSO proved a good initial method for two reasons: relatively easy interpretability and, more importantly, it handles high dimensional and sparse data well. The LASSO algorithm was implemented in R using the package `glmnet` (Friedman et al., 2010) and `tidymodels` framework (Kuhn & Wickham, 2020), where the error regularization penalty was chosen from model training via cross-validation.

In addition, we used gradient boosted trees using the popular eXtreme Gradient Boosting (XGBoost) method (Chen & Guestrin, 2016) as our second model. XGBoost has the benefits of being able to scale well with large, high-dimensional, and sparse data, as well as extremely fast implementation. In general, gradient boosting goes through cycles to iteratively add models into an ensemble. The first step is to initialize the ensemble by creating a single (naive) model. Then the cycle begins, as the current ensemble is utilized to make predictions, which are then used to compute a loss function. Based on this loss function result, a new model is fitted and added to the ensemble. In particular, model parameters are determined using gradient descent so that adding the new model to the ensemble will produce a lower loss. Finally, this new model is added to the ensemble, and then the process is repeated. The XGBoost method was implemented in R using the package `xgboost` (Chen et al., 2021) and `tidymodels` framework, where tuning via cross-validation was used to determine the best set of model hyperparameters.

3.3 Model Evaluation

For both LASSO and XGBoost models, we first partitioned the data using a 70-30 train-test split ratio. Furthermore, a 5-fold cross-validation was implemented to train each model; and

to evaluate the performance of the classifiers, we used the area under the receiver operating characteristic curve (AUC) as our evaluation metric.

Figure 5 shows the ROC curve for our LASSO and XGBoost classification performances on the test data set. We found that a penalty of 0.000339 gave us the best LASSO model in terms of AUC. We used this LASSO model to get the predicted probability values for property appeal on the test set and achieved an AUC value of 0.752.

For the XGBoost algorithm, we considered various hyperparameters associated with model complexity, randomness, and step size as tuning parameters. In terms of model complexity, the optimal combination of parameters consisted of 411 trees, 5 minimum node data points required for further splitting, and a subsample ratio of 0.5. As for randomness, the model selected 13 as the optimal number of randomly sampled predictors. Finally, we obtained a learning rate of 0.0184 which determines the step size. The corresponding AUC value for our XGBoost model on the test set is 0.771.




Figure 5: Receiver operating characteristic curve comparing LASSO and XGBoost classification performances on the test set.

3.4 Other Models

After fitting two machine learning models, we wanted to look into other ways of improving our predictions. One particular problem we have with our dataset is that there is a class imbalance issue with our response variable, the appeal status in 2020. Specifically, only 3.6% of observations had an appeal status of “yes” and for the remaining 96.4%, there was not an appeal (see Table 1). In order to deal with imbalance classification and possibly improve our performance, we considered SMOTE (Chawla et al., 2002), an oversampling and data augmentation technique, short for Synthetic Minority Oversampling Technique. SMOTE works by selecting candidates that are close to the minority class in the feature space. In particular, the minority class is oversampled by taking each observation and introducing synthetic examples along the line segments joining any or all of their k nearest neighbors.

We applied the SMOTE technique via the `themis` package (Hvitfeldt, 2021) in R. To generate the new examples of the minority class, we used a default value of 5 nearest neighbors and a majority-to-minority oversampling ratio of 0.2. This ratio means the minority levels would have $1/5$ as many rows as the majority level as a result of oversampling. After processing the data, we fed this less imbalanced data into both LASSO and XGBoost models. However, the AUC values did not improve for each of the methods. Moreover, further examination led us to the conclusion that the SMOTE models suffered from overfitting, failing to produce better results than the previous two trained models.

3.5 Feature Importance

After fitting the classification models, we examined the importance of the features in the data to the model’s prediction using the SHAP (Lundberg & Lee, 2017) method for model explainability. The name SHAP comes from Shapley values (Shapley, 1951), an idea from game theory which describes the average marginal contribution of a player across all possible coalitions in a game. In a prediction setting, a game can represent a prediction task for an

observation. The players in a game are feature values associated with each observation that collaborate to receive a gain or payout, or predict a certain value. The gain or payout is the difference between the predicted value for a single observation and the average prediction for all observations.

Since the XGBoost method gave us the best classification results, we decided to look into variable importance scores for this model. Figure 6 is a bar graph displaying the top 15 features by importance from the XGBoost model. In addition, Figure 7 is a beeswarm summary plot showing how the top 15 features in the data impact the XGBoost model's output. The four features contributing the most to the prediction of property appeal status were longitude, finished area, whether the house belongs to zip code 53212, and the year built. Each of these features changes the predicted absolute appeal probability on average by at least 7.5 percentage points, as illustrated by Figure 6. In addition, the SHAP beeswarm summary plot indicates that high feature values for longitude, finished area, zip code 53212, and year built are associated with a higher probability of appeal. In context, houses that are more likely to appeal are those located further to the east, with higher total finished area, belonging to zip code 53212, and were built in a more recent year.

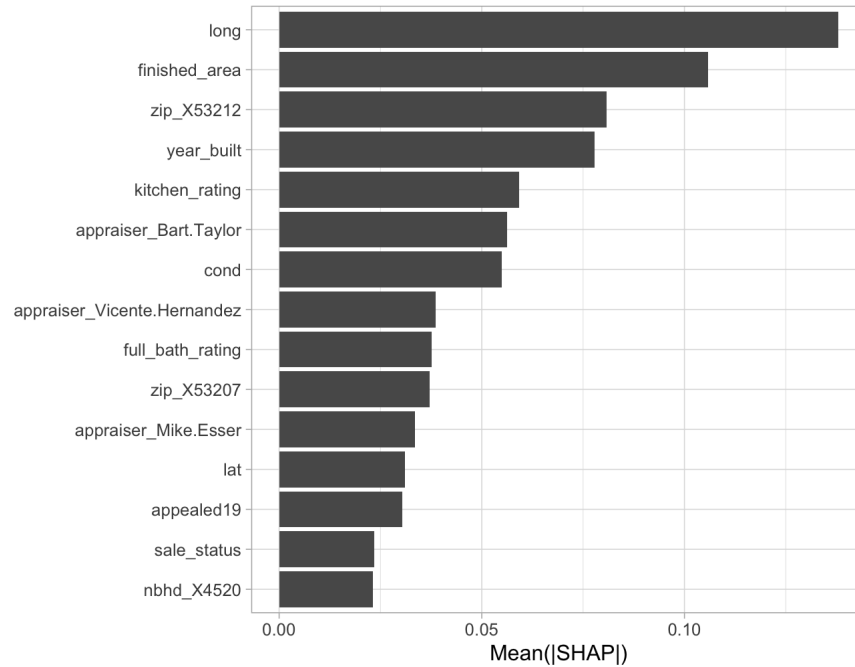


Figure 6: Bar graph of top 15 features by importance for the XGBoost model. SHAP feature importance was measured as the mean absolute Shapley values.




Figure 7: Beeswarm plot of top 15 features by importance for the XGBoost model. The summary plot combines feature importance with feature effects, where each dot represents a Shapley value for a feature and an instance. The color represents feature value from low to high.

4 Conclusion and Discussion

This report investigated trends and proposed methodologies for predicting if a particular property is likely to appeal its valuation assessment. After considering two supervised machine learning methods, LASSO logistic regression and XGBoost, it was found that the XGBoost classification technique performs the best with respect to AUC values. Moreover, it was determined from our XGBoost model that the variables with the greatest contributions to the prediction of property assessment appeals are location (longitude and zip code 53212), year built, and finished area.

Any future application of this methodology would be wise to incorporate the context of previous homeownership appeals. In Chicago, only two hours south of Milwaukee, housing appeals are widely associated with systematic race-based price devaluation (Moore, 2021), thus adding and prioritizing a new motive of recouping lost property value. However, configuring the model to suit more generalized data dating back to before the 2008 housing crisis for a longitudinal analysis could prove very beneficial when attempting to predict long-term housing market trends.

Supplementary Material

All materials related to this manuscript are publicly available on GitHub at <https://github.com/qntkhvn/milwaukee-appeal>. In addition, source code for a Shiny app for visualizing the predicted probabilities of appeal for properties in the city of Milwaukee can be found in the GitHub repository. We chose not to deploy the Shiny app to an open source Shiny server due to sensitive information regarding property location, address, and appraiser.

References

- Cambon, J., Hernangómez, D., Belanger, C., & Possenriede, D. (2021). Tidygeocoder: An r package for geocoding. *Journal of Open Source Software*, 6(65), 3544. <https://doi.org/10.21105/joss.03544>
- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *J. Artif. Int. Res.*, 16(1), 321–357.
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Chen, T., He, T., Benesty, M., Khotilovich, V., Tang, Y., Cho, H., Chen, K., Mitchell, R., Cano, I., Zhou, T., Li, M., Xie, J., Lin, M., Geng, Y., & Li, Y. (2021). *Xgboost: Extreme gradient boosting*. <https://CRAN.R-project.org/package=xgboost>
- Friedman, J., Hastie, T., & Tibshirani, R. (2010). Regularization paths for generalized linear models via coordinate descent. *Journal of Statistical Software*, 33(1), 1–22. <https://www.jstatsoft.org/v33/i01/>
- Hvitfeldt, E. (2021). *Themis: Extra recipes steps for dealing with unbalanced data*. <https://CRAN.R-project.org/package=themis>
- Kuhn, M., & Wickham, H. (2020). *Tidymodels: A collection of packages for modeling and machine learning using tidyverse principles*. <https://www.tidymodels.org>
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Proceedings of the 31st International Conference on Neural Information Processing Systems*, 4768–4777.
- Moore, N. (2021). *Racial inequality in how chicago-area homes are valued is increasing*. <https://www.wbez.org/stories/racial-inequality-in-how-chicago-area-homes-are-valued-is-increasing/241643ab-6cba-4646-9660-af4f26acc6d8>
- R Core Team. (2021). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. <https://www.R-project.org/>

- Shapley, L. S. (1951). *Notes on the n -person game - I: Characteristic-point solutions of the four-person game*. RAND Corporation. <https://doi.org/10.7249/RM0656>
- Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society (Series B)*, 58, 267–288.

Individual Contributions

Together we brainstormed and decided the analysis strategy, questions for the client, wrote and edited the two presentations and final report. Adam and Anthony researched and found literature relevant to the topic (domain knowledge) and performed exploratory data analysis. Quang wrote code for the predictive modeling pipeline.